

Agentic AI Course - Day 7

Krish Naik - 10h course | Video 4:00-4:40

RAG Fundamentals - Loaders & Chunking

RAG (Retrieval-Augmented Generation) lets an LLM answer from YOUR documents. Step one: load documents and split them into chunks small enough to retrieve and fit in the prompt.

Key concepts

- Document loaders: read PDFs, web pages, text, etc. into Document objects.
- Chunking: split long text into overlapping pieces (e.g. 1000 chars, 200 overlap).
- Why chunk: embeddings and context windows work better on focused pieces.
- RecursiveCharacterTextSplitter is the common, robust splitter.

Code example

```
from langchain_text_splitters import RecursiveCharacterTextSplitter

splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=200)
chunks = splitter.split_text(long_text)
print(len(chunks))
```

Today's tasks (1.5h)

- ☐ Load a PDF or webpage
- ☐ Split into chunks with overlap
- ☐ Inspect a few chunks

Resources

Video: youtu.be/rV3HJ4LEZ7k (start at 4:00-4:40)

Code: github.com/krishnaik06/RAG-Tutorials

(Full clickable links are in the app: Goals > this day > Study block.)