

Retrievers & the Full RAG Pipeline

A retriever fetches relevant chunks; you stuff them into the prompt so the LLM answers grounded in your data. Today you connect retrieval + generation end to end.

Key concepts

- Retriever: `vs.as_retriever()` returns top matching chunks for a query.
- Prompt: include the retrieved context + the question.
- Chain: retriever -> prompt -> model -> parser (a RAG chain).
- Always keep sources so answers are verifiable.

Code example

```
retriever = vs.as_retriever(search_kwargs={'k':3})
rag = ({'context': retriever, 'question': lambda x: x} | prompt | llm | StrOutputParser())
print(rag.invoke('What does the doc say about X?'))
```

Today's tasks (1.5h)

- [] Turn the store into a retriever
- [] Build a RAG chain with context+question
- [] Ask questions and check the sources

Resources

Video: youtu.be/rV3HJ4LEZ7k (start at 5:20-6:00)

Code: github.com/krishnaik06/RAG-Tutorials

(Full clickable links are in the app: Goals > this day > Study block.)