

Agentic AI Course - Day 12

Krish Naik - 10h course | Video 7:20-8:00

Guardrails

Guardrails keep LLM apps safe and on-topic: validating inputs, filtering outputs, blocking unsafe content, and enforcing formats. Essential before production.

Key concepts

- Input guardrails: reject prompt-injection, off-topic, or unsafe requests.
- Output guardrails: check the response for PII, toxicity, or format errors.
- Structured validation: enforce JSON schemas / Pydantic on outputs.
- Fail closed: on a guardrail hit, return a safe fallback, not the raw output.

Code example

```
# Validate output against a schema before returning it
from pydantic import BaseModel
class Ans(BaseModel):
    answer: str
    safe: bool
```

Today's tasks (1.5h)

- [] Add an input check for off-topic requests
- [] Validate the output JSON with Pydantic
- [] Return a safe fallback on failure

Resources

Video: youtu.be/rV3HJ4LEZ7k (start at 7:20-8:00)

Code: github.com/krishnaik06/Langchain-V1-Crash-Course

(Full clickable links are in the app: Goals > this day > Study block.)