

Agentic AI Course - Day 13

Krish Naik - 10h course | Video 8:00-8:40

LLM Evaluation

You cannot improve what you do not measure. Evals score your RAG/agent on correctness, faithfulness (no hallucination), and relevance - automatically.

Key concepts

- Faithfulness: is the answer grounded in the retrieved context?
- Answer relevance: does it actually address the question?
- Context precision/recall: did retrieval fetch the right chunks?
- LLM-as-judge: use a model to score outputs against a rubric or ground truth.

Code example

```
# Build a small eval set: [{question, ground_truth}], run the RAG,  
# then score faithfulness/relevance (e.g. with RAGAS or an LLM judge).
```

Today's tasks (1.5h)

- [] Create 5-10 Q/A eval examples
- [] Run your RAG on them
- [] Score faithfulness & relevance; note weak spots

Resources

Video: youtu.be/rV3HJ4LEZ7k (start at 8:00-8:40)

Code: github.com/krishnaik06/RAG-Tutorials

(Full clickable links are in the app: Goals > this day > Study block.)